

Jacob Daniel R

jacobdanielr82@gmail.com | +91 6382243918

[LinkedIn](#) | [Portfolio](#) | [GitHub](#)

Data Engineer focused on building end-to-end data pipelines and cloud-based data platforms. Experienced in SQL, Python, PySpark, and data warehousing. Strong foundation in ETL design, data architecture, and transforming raw data into reliable, high-quality datasets that enable data-driven workflows.

Key Skills

Programming: SQL, Python, PySpark

Data Engineering: ETL/ELT Pipelines, Data Warehousing, Data Modeling, Big Data

Cloud & Tools: Azure Databricks, Power BI, Git/GitHub, Snowflake, AWS, Airflow, Spark, Dbt

Professional Experience

AtliQ Technologies

Remote

Data Engineer (Internship)

Jun 2026 - Aug 2026

- Architected a serverless AWS pipeline (S3, Glue, Lambda, PySpark) validating 1,000+ Finance and Sales records, segregating valid and rejected data with Parquet optimization
- Extended the pipeline into a Silver-layer transformation using Azure Databricks and Delta Lake, standardizing multi-source datasets
- Engineered an event-driven ingestion system for a multi-region financial platform, consolidating 3 Lambda functions into one solution
- Enforced automated validation achieving 97–98% data accuracy at production scale, with schema evolution and real-time alerting

Projects

Automated EOD Securities Pricing Platform - Securities Market Analytics [Project Link](#)

[Apache Airflow / Snowflake / AWS / Python / SQL]

Jul 2026 - Aug 2026

- Orchestrated 2 ingestion modes for historical backfills and daily incremental EOD pricing through Airflow, S3, and Snowflake.
- Streamlined 8 Snowflake tasks with MERGE-based deduplication, validation, and rejected-record handling across RAW and CORE.
- Modeled 3 dimensions and 1 fact table to structure reusable Snowflake datasets for securities pricing analytics.
- Published 6 Snowflake views powering Power BI insights across equity performance, sector liquidity, ETF activity, and watchlists

Azure Databricks Retail Lakehouse - Retail Analytics [Project Link](#)

[Azure Databricks / PySpark / Data Warehousing / ETL Pipelines / SQL]

May 2026 - Jun 2026

- Constructed a Delta Lakehouse in Azure Databricks containing 23 Delta tables processing 1M+ retail records.
- Executed data cleansing, schema evolution, and Delta MERGE to produce trusted analytics-ready datasets.
- Coordinated 4 Databricks workflows and set up Unity Catalog with role-based access control.
- Supplied 8 analytics-ready tables (5 dimensions, 3 facts) powering a Power BI dashboard.

AWS Support Analytics Pipeline - Customer Support Analytics [Project Link](#)

[Amazon S3 / AWS Lambda / AWS Glue / Amazon Redshift / ETL Pipelines]

Apr 2026 - May 2026

- Consolidated 2 data sources (MySQL & logs) into an automated AWS ETL pipeline featuring incremental ingestion.
- Built an event-driven ETL workflow using 5+ AWS services with S3-triggered automation and Parquet storage.
- Validated datasets with Athena and enabled incremental Redshift loading for consistent warehouse updates.
- Minimized manual pipeline execution by ~2 hrs/week by automating transformation, loading, and reporting.

Education

Anna University

Tamil Nadu, India

Bachelor of Engineering - Computer Science and Engineering

Aug 2022 - May 2026

CGPA: 8.38